

Group Details:

- Name: Yonah Huppert ID: 328949813
- Name: Ben Towler ID: 345576433
- Name: Hallel Azulai ID: 212667497
- Name: Ido Zchut ID: 206679029

> Helper Functions and Imports

[Show code](#)

✓ Introduction to Data Science - Lab #2

Exploratory Data Analysis

✓ Case Study: Rental Listings in Jerusalem

In this lab we will practice our exploratory data analysis skills using real data!

We will explore data of rental pricings in Jersualem. The dataset consists of listings published in <https://www.komo.co.il/> during the summer of 2022.

We will use two python packages for visualizing the data: `matplotlib` (and specifically its submodule `pypplot` imported here as `plt`) and `seaborn` (imported as `sns`). Seaborn is a package that "wraps" `matplotlib` and introduces more convenient functions for quickly creating standard visualizations based on dataframes.

Please **briefly** go over this [quick start guide](#) to `matplotlib`, the [first](#) `seaborn` introduction page until the "Multivariate views on complex datasets" section (not included), and the [second](#) introduction page until the "Combining multiple views on the data" section.

✓ Loading the dataset

```
#@title Loading the dataset
rent_df = load_df(RENT_ID)[['propertyID', 'neighborhood', 'monthlyRate', 'mefarsem', 'rooms', 'floor', 'area', 'entry', 'description', 'numFloor:
rent_df = rent_df.drop_duplicates(subset='propertyID').reset_index(drop=True)
rent_df_backup_for_exercise = rent_df.copy()
clean_df_area_filtered = None
clean_df = None
```

Let's print a random sample:

```
np.random.seed(2)
rent_df.sample(5)
```



	propertyID	neighborhood	monthlyRate	mefarsem	rooms	floor	area	entry
403	3981729	גבעת מרדכי	4500.0	private	3.0	6.0	62.0	10/08/2022
457	3991612	קרית חיים	3000.0	private	3.5	3.0	NaN	NaN

And print some summary statistics:

```
rent_df.describe(include='all')
```

	propertyID	neighborhood	monthlyRate	mefarsem	rooms	floor	area	entry	description	numFloors
count	6.120000e+02	612	612.000000	612	612.000000	611.000000	295.000000	292	596	610.000000
unique	NaN	54	NaN	2	NaN	NaN	NaN	17	578	NaN
top	NaN	קרית יובל	NaN	private	NaN	NaN	NaN	10/08/2022	להשכרה, דירה, קומה ראשונה, בירושלים	NaN
freq	NaN	66	NaN	600	NaN	NaN	NaN	259	8	NaN
mean	3.981582e+06	NaN	4717.393791	NaN	2.927288	1.916530	87.664407	NaN	NaN	3.908197
std	6.525543e+04	NaN	2195.215139	NaN	1.007350	1.581006	277.004591	NaN	NaN	1.978065
min	2.494041e+06	NaN	0.000000	NaN	1.000000	-2.000000	1.000000	NaN	NaN	1.000000
25%	3.981694e+06	NaN	3500.000000	NaN	2.000000	1.000000	42.000000	NaN	NaN	3.000000
50%	3.987901e+06	NaN	4400.000000	NaN	3.000000	2.000000	60.000000	NaN	NaN	4.000000
75%	3.992605e+06	NaN	5800.000000	NaN	3.500000	2.000000	85.000000	NaN	NaN	4.000000

The variables we will focus on are:


1. neighborhood: The hebrew name of the neighborhood in jerusalem where the listing is located
2. monthlyRate: The monthly rate (שכר דירה) in shekels
3. rooms: The number of rooms in the apartment
4. floor: The floor in which the apartment is located
5. area: The area of the apartment in squared meters
6. numFloors: The total number of floors in the building

What is the distribution of prices in this dataset?

Q: Plot a histogram with 30 bins of the monthly rates in this dataset:


✓ Solution 1

```
# @title Solution 1
plt.figure(figsize=(8,6))
sns.histplot(rent_df['monthlyRate'], bins=30)
plt.xlabel("Monthly Rate (Shekels)");
```



✓ Solution 2

```
# @title Solution 2
rent_df["monthlyRate"].hist(bins=30, figsize=(8,6))
plt.xlabel("Monthly Rate (Shekels)");
```



We see that the prices distribution peaks around ~3500 Shekels and that it is right skewed, as there are some very expensive apartments. We can also see a peak at zero which makes sense as sometimes listings do not include a price. We would want to filter those out when we analyze prices later on.

Q: Print the number of listings that have no monthly rate:

Solution

```
# @title Solution
print("Number of apartments without a price: ", rent_df['monthlyRate'].value_counts()[0].round(3))
```

Number of apartments without a price: 25

We want to remove those listings, but we don't want to lose these entries, as we might want to know how many and what type of outliers we originally removed. So we create another dataframe that has the listings we removed and the reason for removal.

```
outlier_df = pd.DataFrame(columns=rent_df.columns.to_list()+['reason']) # will save the outliers

outliers = rent_df[rent_df['monthlyRate'] <= 0].reset_index(drop=True)
outliers['reason'] = "monthlyRate <= 0"
outlier_df = pd.concat([outlier_df, outliers], axis=0, ignore_index=True).drop_duplicates().reset_index(drop=True)
outlier_df.tail()
```

	propertyID	neighborhood	monthlyRate	mefarsem	rooms	floor	area	entry	description	numFloors	reason
20	3983978	קריית משה	0.0	private	4.0	3.0	100.0	10/08/2022	דירת 4 חדרים - במצב מצוין - שמורה ביותר כולל ...	5.0	monthlyRate <= 0
21	3985184	נווה יעקב	0.0	private	4.0	1.0	68.0	10/08/2022	דירה במצב שמור מאד. ממוזגת, 2 חדרי שירותים (אח	2.0	monthlyRate <= 0
22	3952750	חולון	0.0	private	5.5	1.0	180.0	10/08/2022	להשכרה דירה מפוארת, רחבת הים מרחק	4.0	monthlyRate

We will now remove those listings and save the result to a new variable `clean_df`:


```
clean_df = rent_df[rent_df['monthlyRate'] > 0].reset_index(drop=True)
```

What is the distribution of the number of rooms?

Q: Use `sns.countplot` to compare the counts of listings with different numbers of rooms. Plot all bars in the same [color](#) of your choice.

✓ Solution


```
# @title Solution
if clean_df is None:
    print("Can't run until 'clean_df' is created!")
else:
    sns.countplot(x='rooms', data=clean_df, color='tab:blue')
    plt.xlabel("# of rooms");
```



The distribution peaks at three rooms and we also see that "half rooms" are less common.

Can we see an association between apartment area and price?

```
if clean_df is None:
    print("Can't run until 'clean_df' is created!")
else:
    plt.figure(figsize=(12,8))
    sns.scatterplot(x='area', y='monthlyRate', data=clean_df)
    plt.ylabel("Monthly Rate (Shekels)")
    plt.xlabel("Area ($m^2$)");
```



We see clear outliers here! We know that area is measured in squared meters and it is unlikely that there are any apartments of $\sim 1000m^2$.

Let's look at those samples to see if we can understand what happened there:

```
if clean_df is None:
    print("Can't run until 'clean_df' is created!")
else:
    display(clean_df.sort_values('area', ascending=False).head(4))
```

	propertyID	neighborhood	monthlyRate	mefarsem	rooms	floor	area	entry	description	numFloors
185	3964340	תלפיות	4400.0	private	2.0	2.0	1234.0	10/08/2022	NaN	3.0
543	3956561	זכרון משה	6600.0	private	3.5	3.0	1000.0	01/07/2022	דירה מהממת בלב ירושלים. צמודה לרכבת הקלה- תחנת...	3.0
576	3974914	תלפיות	17000.0	private	5.0	3.0	220.0	10/08/2022	דירת 5 חדרים ענקית ומהממת, בבנין בוטיק ויחודי	4.0

And inspect the description of one of those listings:

```
if clean_df is None:
    print("Can't run until 'clean_df' is created!")
else:
    display(clean_df.at[543, 'description'])
```

דירה מהממת בלב ירושלים. צמודה לרכבת הקלה- תחנת הדוידקה. 3 חדרים ענקיים ולכל חדר מרפסת גדולה. חלל כניסה עם פינת ישיבה. מתאימה מאוד ל- 3 שותפים'

Clearly not a 1000 m² apartment...


Q: Save a new dataframe named `clean_df_area_filtered` with all listings with area smaller than 800 m². Again, add the removed outliers to the `outliers_df` dataframe.

Plot again the scatter of area vs. monthly rate after removing the outliers.

✓ Solution

```
# @title Solution
if clean_df is None:
    print("Can't run until 'clean_df' is created!")
elif outlier_df is None:
    print("Can't run until 'outlier_df' is created!")
else:
    # save outliers
    outliers = clean_df[clean_df['area'] >= 800].reset_index(drop=True)
    outliers['reason'] = "'area' >= 800"
    outlier_df = pd.concat([outlier_df, outliers], axis=0, ignore_index=True).drop_duplicates().reset_index(drop=True)

    # remove the outliers from the dataset
    clean_df_area_filtered = clean_df[clean_df['area'] < 800].reset_index(drop=True)
    plt.figure(figsize=(12,8))
    sns.scatterplot(x='area', y='monthlyRate', data=clean_df_area_filtered)
    plt.xlabel("Area ($m^2$)")
    plt.ylabel("Monthly Rate (Shekels)");
```



Again, we see some strange behavior of apartments with almost zero area but with a high monthly rate. Let's check them out:

We start with all apartments with an area between 0 to 25 m^2 :

```
# Show all apartments with area between 0 and 25
clean_df_area_filtered[clean_df_area_filtered['area'].between(0,25)]
```

	propertyID	neighborhood	monthlyRate	mefarsem	rooms	floor	area	entry	description	numFloors	
0	3994505	קריית יובל	2000.0	private	1.0	2.0	13.0	10/08/2022	יחידת דיור להשכרה ברחוב הראשי של קריית יובל, ה	2.0	
1	3981298	רחביה	2450.0	private	1.0	1.0	17.0	10/08/2022	דירת יחיד 17 מטר כולל מרפסת קטנה	3.0	
3	3993997	בית וגן	2100.0	private	1.0	0.0	15.0	10/08/2022	דירת חדר, כ-15 מ"ר, במיקום מרכזי אך משופץ	3.0	
5	3993552	הר נוף	2000.0	private	1.0	0.0	20.0	10/08/2022	יחידה משופצת ליחיד או למשרד , מיקום מצוין	4.0	
6	3972039	גבעת שאול	2700.0	private	1.0	0.0	22.0	10/08/2022	דירת חדר כחדשה , כניסה נפרדת ללא וועד בית , מו	1.0	
7	3988096	המושבה הגרמנית	2500.0	private	1.0	0.0	18.0	10/08/2022	רלוונטי לנשים בלבד. ללא עישון. ללא חיות מחמד ...	1.0	
8	3992809	נחלאות	3200.0	private	1.0	2.0	25.0	10/08/2022	הדירה המגניבה בנחלאות, מתפנה אחרי תקופה ארוכה	2.0	
10	3983516	הגבעה הצרפתית	2000.0	private	1.0	2.0	20.0	10/08/2022	דירת חדר קטנה, מסוגננת ונחמדה, מתאימה ליחיד בל	13.0	
12	3987706	נחלאות	2800.0	private	1.0	0.0	20.0	10/08/2022	להשכרה, דירה, בירושלים ברחוב חצור בנחלאות. דיר	2.0	
13	3991842	קטמון הישנה	2500.0	private	1.0	1.0	20.0	10/08/2022	דירת חדר ליחיד באזור יפיפה ושקט בקטמון הישנה	4.0	
14	3992479	קריית יובל	2400.0	private	1.0	1.0	20.0	10/08/2022	דירת חדר חמודה עם גינה קטנה משותפת, מתאים ליחיד	1.0	
15	3974372	תל ארזה	2500.0	private	1.0	0.0	11.0	10/08/2022	להשכרה, דירת חדר, בירושלים	4.0	
19	3985106	קטמונים	2300.0	private	1.0	0.0	20.0	10/08/2022	דירת חדר חמודה עם חצר במיקום מרכזי. כמה דק' ה	3.0	
21	3985295	מוסררה	2600.0	private	1.5	0.0	25.0	10/08/2022	אזור חרדי. מיקום מרכזי קרוב להכול. מרוהטת מלא	4.0	
27	3982008	קריית יובל	2800.0	private	1.5	0.0	25.0	10/08/2022	הדירה נמצאת ברחוב שקט 7 דקות הליכה מעין כרם יש	3.0	

Some make sense and others do not. Let's focus on the expensive ones (between 5,000 and 10,000 shekels):

```
# Show all apartments with area between 0 and 25 that also have a price between 5000 and 10000
if clean_df_area_filtered is None:
    print("Can't run until 'clean_df_area_filtered' is created!")
else:
    display(clean_df_area_filtered[clean_df_area_filtered['area'].between(0,25) & clean_df_area_filtered['monthlyRate'].between(5000, 10000)])
```

	propertyID	neighborhood	monthlyRate	mefarsem	rooms	floor	area	entry	description	numFloors	
197	3984483	ארנונה	6600.0	private	4.0	2.0	1.0	01/09/2022	בשכונת ארנונה, רח' שלום יהודה, דירת 4 חדרים מש	4.0	
234	3985019	פסגת זאב	5500.0	private	4.0	3.0	1.0	10/08/2022	להשכרה 4 חדרים מרווחת עם מרפסת ... סוכה שטופת שמש	4.0	
235	3944204	בית הכרם	7750.0	private	4.0	3.0	10.0	01/09/2022	בבניין חדש, מושקעת, מוארת, מאווררת רח' שקט, רח	5.0	

Those are clearly wrong too... Besides that the relationship between the area and the price seems linear. Let's remove these outliers too:

```
#remove the outliers
if clean_df_area_filtered is None:
    print("Can't run until 'clean_df_area_filtered' is created!")
elif outlier_df is None:
    print("Can't run until 'outlier_df' is created!")
else:
    non_outliers = clean_df_area_filtered['area'] > 10 # get non outliers series of true/false

    # save outliers
    outliers = clean_df_area_filtered[~non_outliers].reset_index(drop=True) # get the outliers
    outliers['reason'] = "'area' <= 10"
    outlier_df = pd.concat([outlier_df, outliers], axis=0, ignore_index=True).drop_duplicates().reset_index(drop=True)

    # remove them
    clean_df_area_filtered = clean_df_area_filtered[non_outliers].reset_index(drop=True)
```


Can we see a different pattern for top floor apartments?

Q: Plot again a scatter of area vs. monthly rate. This time distinguish (by color / marker style or both) between apartments that are in the top floor and the rest of the apartments. (To do that you should create a new column in `clean_df_area_filtered` called `is_top_floor` and set it to 1 if the apartment is in the top floor and 0 otherwise.)

✓ Solution

```
# @title Solution
```

```
if clean_df_area_filtered is None:
    print("Can't run until 'clean_df_area_filtered' is created!")
else:
    clean_df_area_filtered['is_top_floor'] = clean_df_area_filtered['floor'] == clean_df_area_filtered['numFloors']
    plt.figure(figsize=(8,6))
    sns.scatterplot(x='area', y='monthlyRate', data=clean_df_area_filtered, alpha=0.8, hue='is_top_floor', style="is_top_floor");
    plt.xlabel("Area (m^2)")
    plt.ylabel("Monthly Rate (Shekels)");
```



We can take a deeper look on the apartments with the very high monthly rate (to see if those are outliers or not):

```
if clean_df_area_filtered is None:
    print("Can't run until 'clean_df_area_filtered' is created!")
else:
    display(clean_df_area_filtered[clean_df_area_filtered['monthlyRate'] > 11000])
```

	propertyID	neighborhood	monthlyRate	mefarsem	rooms	floor	area	entry	description	numFloors	is_top_floor
198	3956418	רחביה	13000.0	agent	4.0	1.0	100.0	NaN	Beautifully renovated furnished authentic arab...	3.0	False
236	3985051	טלביה	17000.0	private	4.0	4.0	140.0	10/08/2022	4 בקינג דיור רדונס דירת תיוחדרים להשכרה ללא תיו	10.0	False
257	3982363	רחביה	16000.0	private	5.0	2.0	172.0	10/08/2022	הזדמנות נדירה!!! ברחוב רד"ק, בנקודה הכי קרובה	5.0	False
260	3980016	אבו תור	12500.0	private	5.0	0.0	150.0	10/08/2022	דירת דופלקס ענקית 150 מ"ר עם כניסה נפרדת.	4.0	False

We can see some representation of the more expensive neighborhoods of Jerusalem here.. More on the neighborhoods later on!

Is there also a relation between the number of rooms and the listing price?

Q: Create a visualization that compares the distribution of prices for different number of rooms. Your visualization should provide information about central tendency (mean/median/mode) and some information about the distribution of individual values around it (standard


deviation/interquartile range) for each number of rooms. Also, show the real prices of the listings per number of rooms.

✓ Solution

```
# @title Solution
if clean_df_area_filtered is None:
    print("Can't run until 'clean_df_area_filtered' is created!")
else:
    plt.figure(figsize=(10,8))
    sns.boxplot(x='rooms', y='monthlyRate', data=clean_df_area_filtered, color='tab:blue')
    sns.stripplot(x='rooms', y='monthlyRate', alpha=0.4 ,size=4,color='k',data=clean_df_area_filtered)
    plt.xlabel("# of Rooms")
    plt.ylabel("Monthly Rate (Shekels)");

    # Or:
    # plt.figure(figsize=(10,8))
    # sns.barplot(x='rooms', y='monthlyRate', data=clean_df_area_filtered, color='tab:blue', errorbar=None, estimator='median')
    # # Can also use mean but median is more informative in this case as prices are skewed...
    # sns.stripplot(x='rooms', y='monthlyRate', alpha=0.4 ,color='k',data=clean_df_area_filtered)
    # plt.xlabel("# of Rooms")
    # plt.ylabel("Monthly Rate (Shekels)");

    #Violin plot completely fails for very small subsets:
    # plt.figure(figsize=(10,8))
    # sns.violinplot(x='rooms', y='monthlyRate', data=clean_df_area_filtered, color='tab:blue')
    # plt.xlabel("# of Rooms")
    # plt.ylabel("Monthly Rate (Shekels)");
```



Now that we finished pre-processing the data, we can see the state of our outliers VS the data that remains:

```
if outlier_df is None:
    print("Can't run until 'outlier_df' is created!")
else:
    # describe the outlier data
    display(outlier_df.groupby('reason').describe())
    print(f"Proportion removed: {100*len(outlier_df) / (len(outlier_df)+len(clean_df_area_filtered)):.0f} %")
```

22:28 ,30.5.2024

DS4All_Lab_2_Exploratory_Data_Analysis_(EDA) (2).ipynb - Colab

	monthlyRate								rooms					area			numFloors		
	count	mean	std	min	25%	50%	75%	max	count	mean	...	75%	max	count	mean	std			
reason																			
'area' <= 10	5.0	5870.0	1399.821417	4000.0	5500.0	5500.0	6600.0	7750.0	5.0	3.90	...	1.0	10.0	5.0	3.80	1.095445			
'area' >= 800	2.0	5500.0	1555.634919	4400.0	4950.0	5500.0	6050.0	6600.0	2.0	2.75	...	1175.5	1234.0	2.0	3.00	0.000000			
monthlyRate <= 0	25.0	0.0	0.000000	0.0	0.0	0.0	0.0	0.0	25.0	3.28	...	99.0	4554.0	25.0	3.48	2.293469			

3 rows x 40 columns

Submission Exercises

Part 1: Diving deeper into rental prices

We will create a copy of the dataset and work on that. We want to make sure that we do not modify the original dataset.

Part 1 - Create a DataFrame

Show code

Let's go back to the distribution of monthly rental prices in the dataset. Are there interesting trends in the distribution that we missed in the visualizations before?

Use only part1_df for the coding questions in this part



Question 1


Plot 3 different histograms of the monthly prices with 20, 60 and 120 bins respectively, each in a different axis/figure.

```
# Part 1 - Question 1
# Your code goes here:
#hist1
fig = plt.figure(figsize=(8,8))
sns.histplot(x='monthlyRate', data=part1_df, bins=20)
plt.xlabel("Monthly Rate")
plt.ylabel("count");

#hist2
fig = plt.figure(figsize=(7,7))
sns.histplot(x='monthlyRate', data=part1_df, bins=60)
plt.xlabel("Monthly Rate")
plt.ylabel("count");

#hist3
fig = plt.figure(figsize=(10,10))
sns.histplot(x='monthlyRate', data=part1_df, bins=120)
plt.xlabel("Monthly Rate")
plt.ylabel("count");
```






✓ Question 2

For 60 and 120 bins, you can see a repeating pattern of "peaks" and "vallies" in the distribution (mostly in the range between 500 and 7000). Is this pattern due to people rounding the rental prices? Please create a visualization that answers this question. Describe in words how the graph shows what the answer is (Hint: you can use the '%' operator to compute the remainder of dividing values in a pandas Series by a scalar number).

› **extra hint:** please open this cell only after discussing with the course staff the best solution you could come up with

[Show code](#)

```
# Part 1 - Question 2
# Your code goes here:
fig = plt.figure(figsize=(8,8))
part1_df['remainders'] = part1_df['monthlyRate'] % 100
sns.histplot(x='remainders', data=part1_df, bins=20)
plt.xlabel("Monthly Rate Remainders")
plt.ylabel("count");
```




Part 1 Question 2 - textual Answer:

We created a new column in the data that presents the remainder after dividing the price of the apartment by 100. The assumption is that if most of the data has a remainder of zero, it means that the apartment price was likely rounded to the nearest hundred. Our results show this is indeed the case.

✓ Question 3

We expect to see a "drop" in prices frequency near the 5000 Shekels mark due to tax considerations (See [here](#) for an explanation). Create a histogram visualization of the data with the smallest possible bins such that every bin will include exactly one multiplication of 500 (Hint: read the `bins` parameter documentation and what types it accepts). Explain why does this choice of bin size ensures that we will not see rounding effects. Do you see a "drop" around 5000 Shekels? Are there other "drops"?

```
# Part 1 - Question 3
# Your code goes here:
fig = plt.figure(figsize=(20,20))
sns.histplot(x='monthlyRate', data=part1_df, bins = 35)
plt.xlabel("Monthly Rate")
plt.ylabel("count");
#part1_df["monthlyRate"].hist(bins=np.arange(0, part1_df["monthlyRate"].max(), 500))
```



Part 1 Question 3 - textual Answer:

''' This bin choice avoids an exaggerated rounding effect and allows us to see trend at a reasonable rate of increase of price (per 500 shekel). We do see a drop at 5000 shekel, and another significant drop after 7500 and 10000, where it drops almost entirely, then only sporadic instances of higher values appear.

✓ Part 2: Size or number of rooms?

Part 2 - Create a DataFrame for Part 2

```
# @title Part 2 - Create a DataFrame for Part 2
```

```
# Create the dataframe and remove the outliers we found in the intro part:
part2_df = rent_df_backup_for_exercise.copy()
part2_df = part2_df[part2_df['monthlyRate'] > 0].reset_index(drop=True);
part2_df = part2_df[part2_df['area'] < 800].reset_index(drop=True)
part2_df = part2_df[part2_df['area'] > 10].reset_index(drop=True)
```


We saw that both the number of rooms and the area of an apartment are strongly associated with the monthly rate. We now want to check if those are just two perspectives of the same relation (how big is the apartment) or is there something more to it. We will use the cleaned dataframe for this exercise.

Use only `part2_df` for the coding questions in this part

Question 1

Generate a visualization to show that there is a strong association between the number of rooms and the area of the apartment. Explain your choice of plot type and your conclusion from the graph.

```
# Part 2 - Question 1
# Your code goes here:
plt.scatter(part2_df['area'], part2_df['rooms'], s = 12)
plt.xlabel("Area")
plt.ylabel("Rooms");
```



Part 2 Question 1 - textual Answer:

Based on this scatter plot, it is easy to see a trend of more rooms as the size of the apartment increases.

Question 2


Add a new column to the dataframe named "averageRoomSize" with the average room size in the given listing.

```
# Part 2 - Question 2
# Your code goes here:
part2_df['averageRoomSize'] = part2_df['area']/part2_df['rooms']
```

Question 3

Create a plot of the relation between the average room size and the monthly rate.

```
# Part 2 - Question 3
# Your code goes here:
plt.scatter(part2_df['averageRoomSize'], part2_df['monthlyRate'], s = 12)
plt.xlabel("Average Room Size")
plt.ylabel("Monthly Rate");
```



✓ Question 4 - **bonus**

We can see that the variance of the monthly rate increases with the average room size.

Suggest what might be the reason for the increase in the variance and create a visualization to support or refute your suggestion.

```
# Part 2 - Question 4
# Your code goes here:
#
#
```

Part 2 Question 4 - textual Answer:

Write your answer here:

✓ Part 3: Neighborhoods

✓ Part 3 - Function Definitions and DataFrame Creation


```
# @title Part 3 - Function Definitions and DataFrame Creation
```

```
def reverse_string(a):
    return a[::-1]
```

```
socialrank_df = load_df(SOCIORANK_ID)
```

```
neighborhood_ranks = {k: v for k,v in zip(socialrank_df['neighborhood'], socialrank_df['socioEconomicRank'])}
```

```
def get_neighborhood_rank(neighborhood):
    if neighborhood in neighborhood_ranks:
        return neighborhood_ranks[neighborhood]
    else:
        return None
```

```
# Create the dataframe and remove the outliers we found in the intro part:
```

```
part3_df = rent_df_backup_for_exercise.copy()
```

```
part3_df = part3_df[part3_df['monthlyRate'] > 0].reset_index(drop=True);
```

```
part3_df = part3_df[part3_df['area'] < 800].reset_index(drop=True)
```

```
part3_df = part3_df[part3_df['area'] > 10].reset_index(drop=True)
```

```
part3_df["neighborhood_flipped"] = part3_df["neighborhood"].apply(reverse_string) # making the neighborhood names readable
```

We now want to focus on the differences between different neighborhoods in Jerusalem.

Use only part3_df for the coding questions in this part

*Use the "neighborhood_flipped" column for visualizations as seaborn will flip the order of letters in hebrew.

✓ Question 1

Print the number of unique neighborhoods that appear in the dataset.

```
# Part 3 - Question 1
```

```
# Your code goes here:
```

```
print(len(part3_df['neighborhood'].unique()))
```

→ 46

✓ Question 2

Visualize the number of listings per neighborhood in a way that will allow you to easily identify those with the highest count.

```
# Part 3 - Question 2
```

```
# Your code goes here:
```

```
fig = plt.figure(figsize=(50,20))
```

```
sns.histplot(x='neighborhood_flipped', data=part3_df)
```


```
plt.xlabel("Neighborhood")
```

```
plt.ylabel("count");
```

```
plt.xticks(rotation=45, ha='right')
```

```
↳ ([0,
1,
2,
3,
4,
5,
6,
7,
8,
9,
10,
11,
12,
13,
14,
15,
16,
17,
18,
19,
20,
21,
22,
23,
24,
25,
26,
27,
28,
29,
30,
31,
32,
33,
34,
35,
36,
37,
38,
39,
40,
41,
42,
43,
44,
45],
[Text(0, 0, 'לבווי תיירק'),
Text(1, 0, 'היכחר'),
Text(2, 0, 'החלמ'),
Text(3, 0, 'ונו תיב'),
Text(4, 0, 'באז תגספ'),
Text(5, 0, 'פונ רה'),
Text(6, 0, 'לואש תעבג'),
Text(7, 0, 'תינמרגה הבשומה'),
Text(8, 0, 'תואלחנ'),
Text(9, 0, 'העקב'),
Text(10, 0, 'חיתפרצה העבגה'),
Text(11, 0, 'הנשיה וומטק'),
Text(12, 0, 'הזרא לת'),
Text(13, 0, 'פינומטק'),
Text(14, 0, 'השמ תיירק'),
Text(15, 0, 'פינג ריע'),
Text(16, 0, 'הליג'),
Text(17, 0, 'הררסומ'),
Text(18, 0, 'לאומש תיירק'),
Text(19, 0, 'תויפלת'),
Text(20, 0, 'וקסר'),
Text(21, 0, 'ריעה זכרמ'),
Text(22, 0, 'סרכה תיב'),
Text(23, 0, 'הדוהי הנחמ'),
Text(24, 0, 'בקעי הוונ'),
Text(25, 0, 'תינוויה הבשומה'),
Text(26, 0, 'סייח רוקמ'),
Text(27, 0, 'סחנמ תיירק'),
Text(28, 0, 'סואלה תיירק'),
Text(29, 0, 'ביצנה וומרא'),
Text(30, 0, 'רות ובא'),
Text(31, 0, 'תומר'),
Text(32, 0, 'דורב רוקמ'),
Text(33, 0, 'השמ וורכו'),
Text(34, 0, 'לארשי תיב'),
Text(35, 0, 'תרש תמר'),
Text(36, 0, 'המואה תונכשמ'),
Text(37, 0, 'לוכשא תמר'),
Text(38, 0, 'הנונרא'),
Text(39, 0, 'תויפלת וופצ'),
Text(40, 0, 'יכדרמ תעבג'),
Text(41, 0, 'המלש תמר'),
Text(42, 0, 'היבלט'),
Text(43, 0, 'תוינ')])
```

```
Text(44, 0, 'תפ'),
Text(45, 0, 'המוח רה'))]]
```



Question 3 - Heavy-tailed distributions

Print the number of neighborhoods with less than 5 listings and the fraction of their total number of listings out of the total number of listings. Also print the fraction of listings from the 8 most frequent neighborhoods out of the total number of listings.

```
# Part 3 - Question 3
# Your code goes here:
neighborhood = 0
listing = 0
counts = part3_df.groupby('neighborhood')['monthlyRate'].count()
for amount in counts:
    if amount < 5:
        neighborhood += 1
        listing += amount
print("Number of neighborhoods with less than 5 listings: ", neighborhood)
print("Fraction of listings with less than 5 listings out of total listings: ", listing/sum(counts))
print("Fraction of listings from top 8 neighborhoods:", sum(counts.nlargest(8))/sum(counts))
```

```
→ Number of neighborhoods with less than 5 listings: 28
Fraction of listings with less than 5 listings out of total listings: 0.23443223443223443
Fraction of listings from top 8 neighborhoods: 0.5238095238095238
```

Those types of distributions where there are many categories that appear only a few times but together take a large portion of the distribution are called heavy-tailed (or long-tailed) distributions. This is a real issue in many data science applications, since even if we have a large dataset there are still some sub-populations or sub-categories that are not well represented.

Question 4

Create a new filtered dataframe with listings from only the 8 most frequent neighborhoods.

```
# Part 3 - Question 4
# Your code goes here:

counts = part3_df.groupby('neighborhood')['monthlyRate'].count().nlargest(8).sort_values(ascending = False).index
filtered_listings = part3_df[part3_df['neighborhood'].isin(counts)]
```

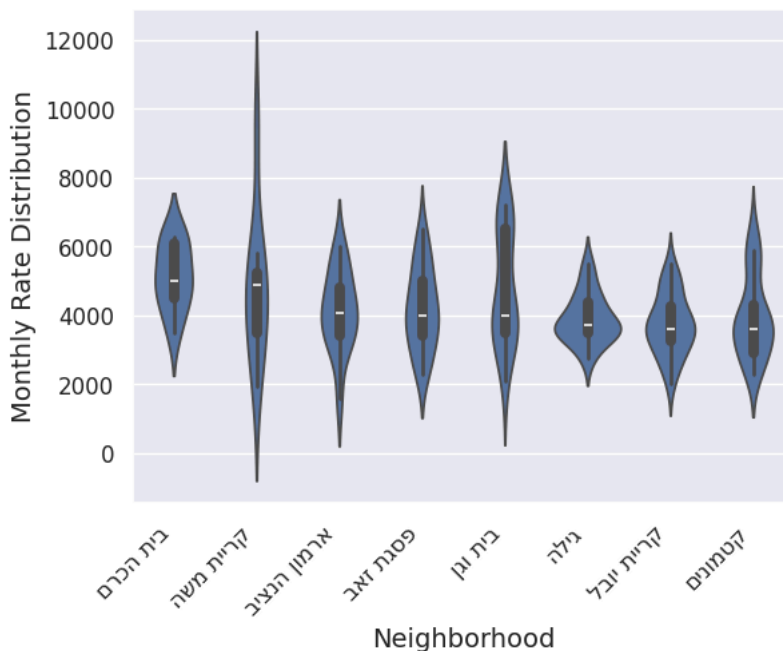
Question 5

Plot a graph to check whether there are different distributions of monthly rates in the eight neighborhoods. Explain your choice for the visualization and your conclusions. Note: Make sure that the neighborhoods are ordered in the plot based on their tendency for higher or lower monthly rates.

Hint: Which is a better descriptor of the central tendency of monthly rates when the distributions are skewed?

```
# Part 3 - Question 5
# Your code goes here:
sns.violinplot(x='neighborhood_flipped', y='monthlyRate', data=filtered_listings, order = filtered_listings.groupby('neighborhood_flipped')
plt.xlabel("Neighborhood")
plt.ylabel("Monthly Rate Distribution")
plt.xticks(rotation=45, ha='right')
```

```
[[0, 1, 2, 3, 4, 5, 6, 7],
 [Text(0, 0, 'מרכה תיב'),
  Text(1, 0, 'השמ תירק'),
  Text(2, 0, 'אומרא וביצנה'),
  Text(3, 0, 'באז תגספ'),
  Text(4, 0, 'תוגו תיב'),
  Text(5, 0, 'הליג'),
  Text(6, 0, 'לבוי תירק'),
  Text(7, 0, 'מינומטק')]]
```



Part 3 Question 5 - textual Answer:

A violin plot allows us to easily compare the distributions of neighborhood prices to each other, centered around the median. It is clear that Beit Hakerem has the largest distribution, and the others can be easily visualized and compared as well.

Question 6

Now that we compared the different distributions of monthly rates between neighborhoods, we can check whether we can explain some of the differences using our common-sense and the data we already have. For example, perhaps different neighborhoods have different distributions of apartment sizes?


Think of a new variable that will allow you to check the relationship between neighborhoods and prices fairly, factoring different apartment sizes out of the equation. Save this measure into the dataframe and create a new visualization to answer the question.

Part 3 - Question 6

Your code goes here:

```
filtered_listings['proportional_price'] = filtered_listings['monthlyRate'] / filtered_listings['area']
```

```
sns.violinplot(x='neighborhood_flipped', y='proportional_price', data=filtered_listings, order = filtered_listings.groupby('neighborhood').xlabel("Neighborhood"))
plt.ylabel("Proportional Monthly Rate Distribution")
plt.xticks(rotation=45, ha='right')
```


 <ipython-input-171-0584c03434bf>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-

```
filtered_listings['proportional_price'] = filtered_listings['monthlyRate'] / filtered_listings['area']
```

```
[[0, 1, 2, 3, 4, 5, 6, 7],
```

```
[Text(0, 0, 'בית וגן'),  
Text(1, 0, 'בית הכרם'),  
Text(2, 0, 'קרית משה'),  
Text(3, 0, 'קטמון'),  
Text(4, 0, 'קרית יואל'),  
Text(5, 0, 'גילה'),  
Text(6, 0, 'ארמון הנציב'),  
Text(7, 0, 'פסגת זאב')]]
```



✓ Part 3 Question 6 - textual Answer:

By dividing the monthly rate by the apartment area, we created a new value that shows the monthly cost of the apartment proportional to its area, to remove its influence on our results. With this new information, the same comparison shows that Bait Vegan has the highest median rent per square meter.

Given the conclusions from the previous steps, we may think that the apartment's neighborhood gives us additional information about the expected monthly rate. But the sample size for most neighborhoods is rather small. So let's examine another way to utilize the location information. Luckily, we also have data about the socio-economic rank of most neighborhoods (between 1 and 10).

✓ Question 7 - bonus

Use again the full dataset (without filtering by neighborhood).

Create an aggregated dataframe where every record represents a neighborhood, with columns for:

1. neighborhood name
2. flipped neighborhood name
3. The number of listings in a neighborhood
4. The median monthly rate for listings in this neighborhood.

Add a column with the neighborhood socio-economic rank to the dataframe (you can use the provided `get_neighborhood_rank` function that takes as an input a neighborhood name and returns its socio-economic rank.) Use this dataframe to visualize the association between socio-economic rank and pricing for all neighborhoods with at least 5 listings. What is your conclusion?

```
# Part 3 - Question 7
# Your code goes here:
#
#
```

Part 3 Question 7 - textual Answer:

Write your answer here:

✓ Part 4: Are private houses more expensive than apartments?

✓ Part 4 - Create a DataFrame and remove outliers for Part 4

```
# @title Part 4 - Create a DataFrame and remove outliers for Part 4
part4_df = rent_df_backup_for_exercise.copy()
part4_df = part4_df[part4_df['monthlyRate'] > 0].reset_index(drop=True);
part4_df = part4_df[part4_df['area'] < 800].reset_index(drop=True)
part4_df = part4_df[part4_df['area'] > 10].reset_index(drop=True)
```

Finally, we want to check if listings in private houses tend to be more expensive than apartments in a building.

Use only `part4_df` for the coding questions in this part

✓ Question 1

The current dataset doesn't include a variable that describes whether a listing is in a building or a private house but this can be inferred from the existing variables. Create a new column named "is_a_house" with value of `True` if a listing is in the first (or zero) floor in a building with only one floor. Print the number of private houses and print the descriptions of three random listings with 'is_a_house' equal to `True`.

```
# Part 4 - Question 1
# Your code goes here:
part4_df['is_a_house'] = (part4_df['floor'] < 2.0 ) & (part4_df['numFloors'] == 1.0)
print(len(part4_df[part4_df['is_a_house'] == True]))

np.random.seed(0)
print(part4_df[part4_df['is_a_house'] == True].sample(3)['description'])
```

```
17
6 דירת חדר כחדשה , כניסה נפרדת ללא וועד בית , מו
34 .בלב נחלאות, דירה מתוקה להשכרה עם כניסה פרטית
48 מקום מדהים ,קומה ראשונה,תחנת אוטובוס,גני ילדים
Name: description, dtype: object
```

✓ Question 2

Create a visualization that compares the **average** monthly rates in houses vs. apartments. Which are more expensive on average?

```
# Part 4 - Question 2
# Your code goes here:
sns.boxplot(x='is_a_house', y='monthlyRate', data=part4_df, meanline=True, showmeans=True)
plt.xlabel("Is it a house?")
plt.ylabel("Monthly Rate")
```